

Deadlock

System Model

- **Computer systems** consist of **resources** (for example CPU, memory, and I/O devices).
- We also consider tools like **locks, and semaphores** to be **resources**.
- Each resource of type R_i has W_i instances.
- When a **process** utilizes a resource the following happens:
 1. The process **requests** the resource.
 2. The process **uses** the resource.
 3. The process **releases** the resource.
 - Processes can request as many resources as they need, it is not restricted. However, the number of resources requested cannot exceed the amount of resources available.
 - Processes may need to wait for resources to become available before their request is granted.
- The operating system uses a **system table** to record the status of **resources** (unallocated / allocated). If a resource is allocated, the process / thread allocating it is also recorded.

Deadlock

- A set of **threads** are said to be **in a deadlocked state** when **every thread** in the set is **waiting for an event** that can be caused only by **another thread in the set**.
- The following is an example of how deadlock can occur:
 1. Process one acquires a lock on resource 1.
 2. Process two acquires a lock on resource 2.
 3. Process wants to acquire a lock on resource two, but must block until it is available.
 4. Process two wants to acquire a lock on resource one, but must block until it is available.
 5. Both processes have entered deadlock.
- A set of **threads** are said to be **in a livelocked state** when **every thread** in the set is **waiting for an event** that can be caused only by **another thread in the set**. Livelock occurs when a thread that the others are waiting for **continuously attempts an action that fails**.

Deadlock Conditions

- **Deadlock** can arise if **four conditions** occur at the same time:
 1. **Mutual exclusion:** Only one process at a time can use the resource.
 2. **Hold and wait:** A process holding at least one resource is waiting to acquire additional resources held by other processes.
 3. **No preemption:** A resource can be released only voluntarily by the process holding it.
 4. **Circular wait:** There exists a set of waiting processes such that P_0 is waiting for P_1 , and P_1 is waiting for P_2 and ... and P_n is waiting for P_0 .
- These conditions do **not guarantee deadlock** will occur, they can **potentially cause deadlock**.

Resource-Allocation Graphs

- A **resource-allocation graph** is a **directed graph**, $G = (V, E)$.
- V is partitioned into two sets:
 1. $P = \{P_1, P_2, \dots, P_n\}$ consisting of all of the processes in the system.
 2. $R = \{R_1, R_2, \dots, R_m\}$ consisting of all of the resource types in the system.
- E consists of two types of edges:
 1. A **request edge** is a directed edge of the form $P_i \rightarrow R_j$.
 2. An **assignment edge** is a directed edge of the form $R_j \rightarrow P_i$.
- If G has a **cycle**, **deadlock** may exist.
- If G has **no cycles**, **deadlock** does not exist.

Methods for Handling Deadlock

- One way to handle deadlocks is to ensure the system will never enter a deadlocked state (prevention / avoidance).
- Another way to handle deadlocks is to allow the system to enter a deadlocked state, and then recover.
- The last way is to ignore the problem, and pretend that deadlocks never occur in the system (this option is used by most systems).